

# Supervised Learning: Classification

Tree-based methods and Support Vector Machines

---

Maarten Cruyff

## Afternoon session

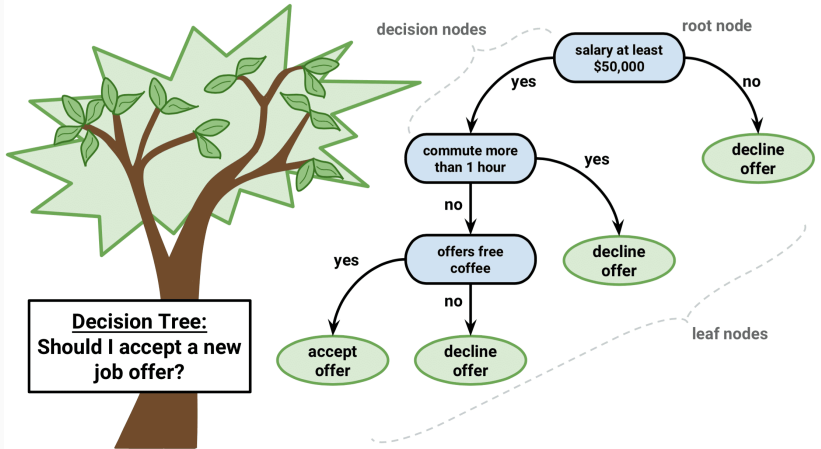
### Trees and SVMs

In this session we look at tree-based methods for classification. A difference with logistic regression is that tree-based methods do not estimate coefficients for the features, but instead partition the feature space in a way that optimizes the predictions. The basic approach is to start with splitting the feature that yields the largest improvement, then find the next feature to split, and so on until the improvement is negligible. Since this approach is very much data-driven, it has high variance. Random forests, bagging, boosting and SVMs are tree-based methods that take measures that protect against high variance, at the cost of interpretability.

### Course materials

- [Lecture sheets](#)
- [R lab](#)
- [R Markdown lab template](#)


# Classification trees



1. Classification trees
2. Pruning
3. Bagging, random forests
4. Boosting
5. Support Vector Machines

# Classification trees

1. Recursive binary splitting algorithm
2. Splits features on basis of node *purity*
  - Gini index
  - deviance



**Figure 1:** Sale of car seats (Yes/No)

# Recursive binary splitting

## Algorithm


1. Divide feature space in non-overlapping, rectangular regions
2. Choose splits that minimize *node impurity* (homogeneity of nodes)
3. Assign region to class with highest mode
4. Stop when node purity no longer increases

Algorithm is top-down and greedy, so

- high variance

# Classification with trees or regression?

Depends on nature of relationship between classes and features



Growing and plotting trees with function `tree()`

```
fit_tree <- tree(formula, data, split = c("deviance", "gini"))  
  
plot(fit_tree)  
text(fit_tree)
```

- minimization of deviance or gini impurity
- `text()` for adding labels to nodes



# Methods to reduce variance:

## 1. Pruning

- cut branches with cross-validation and regularization

## 2. Bagging

- average predictions of bootstrapped trees

## 3. Random forests

- average predictions of decorrelated bootstrapped trees

## 4. Boosting

- weighted combination of weak classifiers (small trees)

# Pruning

---

# Cost-complexity pruning (package tree)

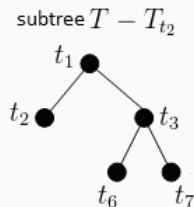
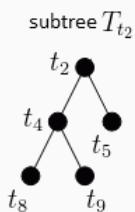
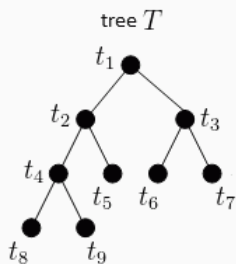
1. Cross-validate tree to find optimal number of nodes
  - based on either deviance or misclassification error rate
2. Prune the tree with optimal number of nodes
  - number of nodes with smallest deviance/misclassification error rate
3. Get predictions of pruned tree on test set
  - `predict.tree` returns either probabilities or classifications

```
cv.tree(fit_tree, method = c("deviance", "misclass"))
```

```
pruned_tree <- prune.tree(fit_tree, best = <number>)
```

```
predict(pruned_tree, newdata, type = "class")
```

# Pruning example

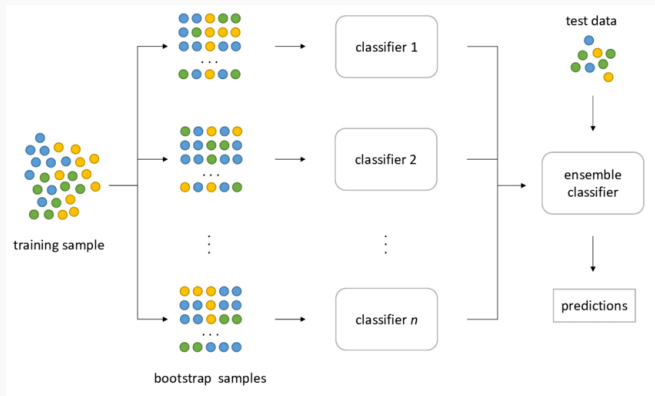


# Bagging, random forests

---

# Algorithm

1. Fit classification trees to  $B$  bootstrap samples
2. Average the predictions
3. Out-Of-Bag (OOB) as estimate validation error



# Bagging vs random forests

## Bagging

- considers all predictors at each split
- best predictors turn up in each tree
- highly correlated trees
- high variance

## Random forests

- considers random sample of predictors at each step
- all predictors get a fair chance
- decorrelated trees
- lower variance

# Out-of-bag error rate

On average  $1/3$  of observations not in bootstrap (Out-Of-Bag)



- OOB cases used to compute validation error
- no need for cross validation
- computationally very efficient



# Variable importance

When averaging trees the tree structure is lost

- how to interpret solution then?
- effect predictors averaged over trees
- visualize with *variable importance* plots



# Package randomForest

Functions for bagging/random forest

```
fit <- randomForest(formula, data,  
                    ntree    = 500,  
                    mtry     = <number predictors at each split>,  
                    importance = TRUE)
```

```
varImpPlot(fit)
```

```
predict(fit, newdata, type = "prob")
```


- mtry: default random forest (bagging total number predictors)
- ntree is tuning parameter (overfitting when too large)
- importance = TRUE for varImpPlot()
- default type yields predicted class

# Boosting

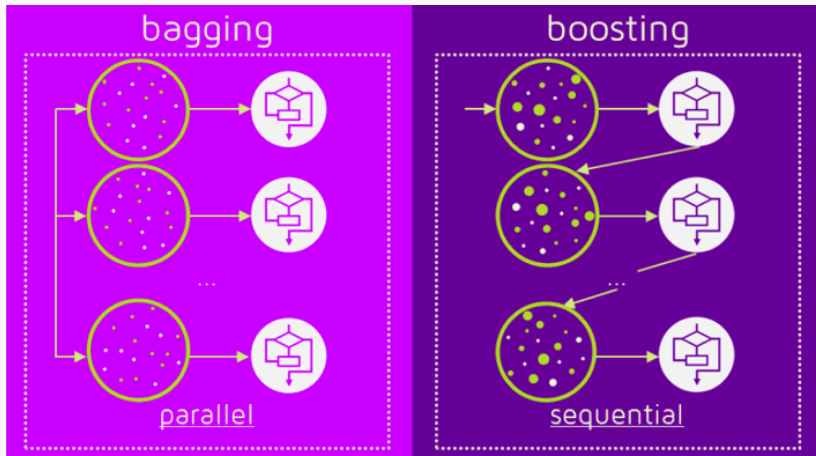
---

# Algorithm

1. Apply a weak classifier (e.g. stump) to training data
2. Increase weights for incorrect classifications, and repeat
3. Classifier is linear combination of weak classifiers



# Boosting vs bagging/random forest



# Boosting with package fastAdaboost

Boosting a single model

- nIter is number of weak classifiers

```
ada <- adaboost(formula, data, nIter)
```

```
predict(ada, newdata)
```

- nIter is number of weak classifiers (overfitting when too large)
- predictions include classes, probabilities and misclassification error

# Boosting with package caret

Determine nIter with cross validation

- use the method adaboost in the function train()
- specify a sequence of values for nIter

```
ada <- train(formula,
             data,
             method = "adaboost",
             trControl = trainControl(method = "cv", number = 5),
             tuneGrid = expand.grid(method = "Adaboost.M1",
                                   nIter = <test sequence>))

predict(ada, newdata, type = "prob")
```

- “Adaboost.M1” restricts search to one of two methods
- default type yields predicted class

# Support Vector Machines (SVM)

---



# SVM for binary classification

## Classifiers using support vectors

### 1. *maximal margin classifier*

- classes perfectly separable by hyperplane

### 2. *support vector classifier*

- allows for non-separable cases

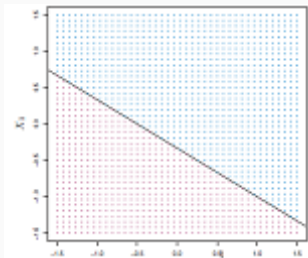
### 3. *support vector machine*

- allows for non-linear boundaries

# Hyperplane

Divides the feature space in two

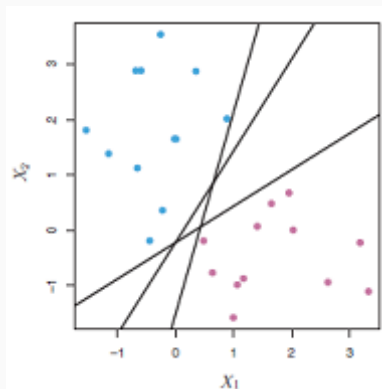
- in two dimensions hyperplane is simply a line



# Separating hyperplane

Perfectly separates the two classes of the outcome variable

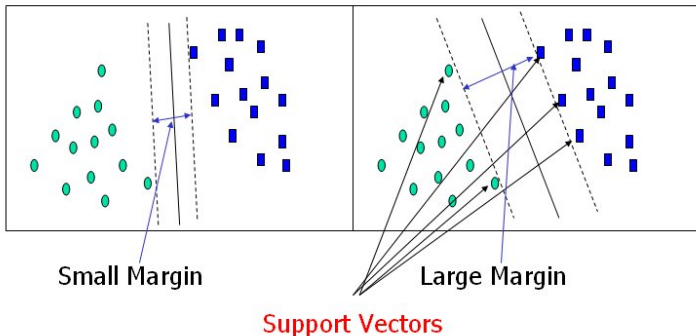
- hyperplane not uniquely identified
- high variance



# Maximal Margin Classifier

Identifies hyperplane by specification of a maximal margin

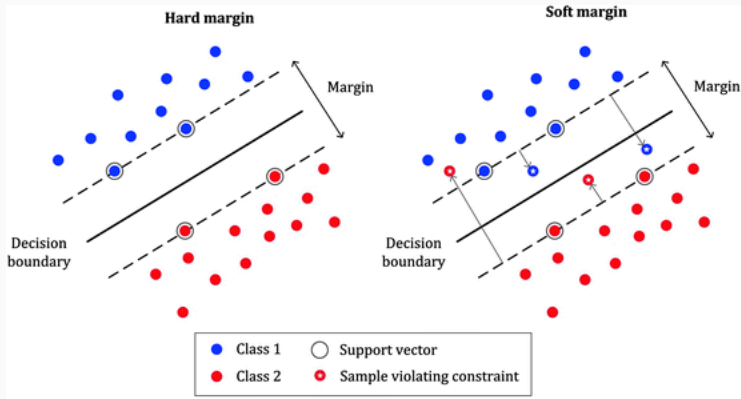
- points on margin are support vectors
- only works if cases are *separable*



# Support Vector Classifier (SVC)

Allows for violations of the margin (soft margin)


- budget for violations is called *cost* ( $C$ )
- cases the wrong side of hyperplane contribute to the cost



# Support Vector Machines (SVM)

Kernels allow for nonlinear hyperplanes, e.g.

- polynomial kernel (left)
- radial kernel (right)



# SVM with package e1071

```
svm_train <- tune(svm, formula, data,  
                 degree = 3, #default  
                 coef0   = 0, #default  
                 cost     = 1, #default  
                 kernel   = c("linear", "polynomial", "radial"),  
                 ranges   = list(cost = <sequence>), etc.)  
  
svm_train$best.model # performance summary  
  
svm_class <- predict(svm_train, newdata, probability = TRUE)  
svm_prob  <- attr(svm_class, "probabilities")
```

- cost, degree and coef0 are tuning parameters
- ranges works similar as tuneGrid()

# SVM classification plot

Compression hyperplane two dimensions

```
plot(svm_train$best.model, data, x1 ~ x2)
```





# Pro's and con's classifiers

## BLR

- robust against outliers but potentially unstable

## LDA

- better stability but sensitive to normality violations

## Tree-based methods

- top-down and greedy, so bias-variance control needed
- boosting considered as one of the best methods

## SVM

- similar to BLR (same loss function), but allow for non-linearity